

Task 2 — Values Audit

Company: AuditAI Compliance Pte. Ltd. (Hypothetical RegTech company)

Q1. Mission, Values, Stage, and Capabilities

Mission Statement:

“We make AI governance legible — to the people who build models, the regulators who oversee them, and the customers whose lives depend on them.”

What the company does:

AuditAI Compliance builds jurisdiction-aware AI governance tooling for financial institutions. Our core product is a compliance monitoring platform that tracks AI model performance, fairness metrics, and regulatory posture across multiple jurisdictions in a single dashboard. We translate regulatory requirements — which differ substantially between the US, EU, Singapore, and elsewhere — into machine-readable monitoring configurations that update automatically when regulations change.

Core values:

- Epistemic honesty: We report what the data actually shows, including when a client's model is drifting or producing disparate outcomes, even when that is uncomfortable.
- Proportionality: We design tools that match the regulatory weight of the problem. We do not gold-plate compliance documentation for its own sake.
- Jurisdictional fidelity: We do not flatten regulatory differences to make our product simpler. If US and EU rules genuinely conflict, our tool surfaces that conflict rather than silently resolving it.

Company stage and size:

AuditAI is a Series A company (~3 years old), with approximately 35 full-time employees: 15 engineers, 8 regulatory specialists, 7 commercial/sales, 5 operations. Headquarters: Singapore. London satellite office. Current ARR: SGD 4.2 million. Targeting profitability at SGD 8-10 million ARR.

What we are good at:

- Translating regulatory text into structured rule engines (ex-MAS, ex-FCA, ex-OCC advisors on staff)

- Quantitative fairness analytics — deep expertise in disparate impact measurement and PSI-based drift detection
- Modular architecture allowing jurisdiction configurations to be updated without changing client-side code

Honest self-assessment of weaknesses:

- Thin coverage of APAC regulations beyond MAS (HKMA and RBI guidance absent from current rule set)
- Monitoring tool only: we detect problems, we do not fix them
- Limited GenAI governance coverage — gap vs. EU AI Act Title III obligations

Q2. Whose Perspective Does the Tool Serve?

Primary client: The Chief Compliance Officer (CCO) or Chief Risk Officer (CRO) of a financial institution deploying AI in credit decisions.

Rationale:

We made a deliberate choice to sell to CCOs rather than to regulators directly. The regulatory SaaS market has a slower sales cycle, is politically complex, and often produces tools used to document compliance rather than improve it. CCOs face genuine liability for model failures and have an incentive to use tools that actually work. We sized the market using McKinsey Global Banking Annual Review 2024 data: AI governance tooling in US/EU is growing at ~18% annually; mid-to-large banks budget SGD 200,000-800,000 annually for model risk management tooling. Our pricing of SGD 120,000-350,000/year is positioned within that range.

Key tensions identified:

- CCO vs. Regulator: The CCO wants a tool that demonstrates compliance. The regulator wants a tool that detects genuine risk. These are not the same thing. A model passing all OCC 2026-13 documentation checks can still produce disparate outcomes that harm consumers.
- CCO vs. Consumer: A credit model maximizing risk-adjusted return may systematically disadvantage protected groups. Our tool will detect this — but the decision of what to do belongs to the client. We acknowledge our tool can be used to document that a problem was known without it being remediated. We believe surfacing the data is better than not surfacing it, but we have no clean resolution to this tension.

Q3. Genuine Risk or Documenting Compliance?

Feature we included that serves genuine risk (not required by regulation):

We included a counterfactual fairness check — a feature that tests whether a model's decision for a specific applicant would change if only their protected attribute (race, gender) were altered, holding all else constant. This is not required by OCC 2026-13 or EU AI Act Article 10. The EU AI Act requires bias testing but does not specify counterfactual methods. We included it because disparate impact ratios (the compliance standard) can be gamed: a model can pass an 80% rule while still making individual decisions directly influenced by protected attributes. Counterfactual checks catch this. They make our product harder to sell because they surface findings that group-level metrics miss.

Feature we chose NOT to include (serves documentation over substance):

We deliberately did not build a regulatory report generator that auto-populates OCC model validation templates or EU Annex IV documentation from our monitoring data. Several competitors offer this. We chose not to because it encourages clients to treat documentation as a substitute for substantive review — the green report problem. A well-formatted Annex IV document produced from a drifting model is worse than no document because it creates false assurance. We lose some deals over this choice. We believe it is the right choice.

Q4. Who Bears the Cost if the Tool Gets It Wrong?

The table below maps each failure mode to the specific harm and the stakeholder who bears it:

Failure Mode	Specific Harm	Primary Victim	Secondary Victim
False Negative: misses subgroup PSI breach	Applicants from affected group denied credit or offered worse terms; harm is diffuse and invisible to victims	Consumers (loan applicants)	CCO — regulatory/reputational liability if later discovered
False Positive: jurisdiction misconfiguration applies EU Article 14 to US model	Wasted manual review time; delayed credit decisions	Financial institution (operational cost)	None — recoverable, no consumer harm
Tool drift: applicant pool demographic shift invalidates baselines	Misleading trend data; genuine fairness problems masked by distributional shift	Both institution and consumers	Regulator — receives unreliable compliance reports
Jurisdictional misconfiguration: EU deployment flagged as US	Client believes they are EU-compliant when they are not; EU AI Act fines up to 30M EUR or 6% global turnover	Financial institution	Consumers in EU — lose AI Act protections

Design response to misconfiguration risk: We have implemented jurisdictional configuration

validation with explicit confirmation dialogs and a jurisdiction audit log. We cannot eliminate misconfiguration risk, but we make it deliberate rather than accidental.